

PROFILE

AI platform and reliability engineer with about eight years of experience delivering governed agent applications, cloud platforms, and user-facing services into production.

Builds and operates systems across agent runtime integration, delivery automation, access control, observability, incident response, and production stabilization.

SELECTED EVIDENCE

AWS EKS platform for governed, self-service AX application delivery	Production reliability for an enterprise AI service used by 30,000 employees	UK rollout, incident response, and stabilization for a global B2C generative AI service
---	--	---

CORE SKILLS

Agent Platform	AWS EKS, agent runtime integration, self-service delivery, automatic domains, existing LLM Gateway API integration
Reliability	Grafana, Prometheus, OpenTelemetry, incident response, rollout stabilization, high availability, auto-scaling
Software & Delivery	Python, Bash, Linux, Kubernetes Operators, ArgoCD, GitHub Actions, GitOps, Helm/Helmfile
Security & Governance	Authentication, authorization, platform security controls, policy enforcement, enterprise governance

EXPERIENCE

Samsung Electronics, AX Development Group	Seoul, South Korea
SRE / AI Platform Engineer	Jan. 2026 - Present
- Design and build an internal AX Agent Platform on AWS EKS, adapting the agent runtime and deployment path to enterprise platform standards. - Implemented self-service AX application delivery with automatic domain issuance, authorized-only access, and platform controls for authentication, authorization, security, and governance. - Enabled AX applications to call the existing enterprise LLM Gateway API through platform-side client compatibility and runtime integration. - Supported a Microsoft-collaborative B2C generative AI service through UK rollout readiness, initial production operations, incident response, and stabilization from Jan. to Jun. 2026.	
Samsung Research	Seoul, South Korea
SRE / AI Platform Engineer	Jul. 2024 - Dec. 2025
- Led production delivery and application reliability for a company-wide AI chatbot serving 30,000 employees on internal Kubernetes clusters. - Owned releases across development, staging, and production while integrating application delivery with security and operational requirements. - Built container delivery pipelines with Docker, Kubernetes, GitHub Actions, ArgoCD, Helm, and Helmfile. - Deployed Grafana, Prometheus, and OpenTelemetry observability for shared service visibility and operational diagnosis.	
Samsung Electronics, Network Division	Seoul, South Korea
SRE & Software Engineer	Jan. 2021 - Jun. 2024
- Designed and delivered containerized applications with Docker, Helm, Kubernetes Operators, and AWS services including EKS, ECS, Route 53, RDS, and IAM. - Built and managed OpenShift, Kubernetes, and EKS environments for reliability, high availability, performance, and auto-scaling. - Designed and developed a Bash-based one-click Kubernetes cluster provisioning system for cloud developers. - Developed a Linux application that mirrors and forwards up to 130,000 packets per second under constrained system resources.	

EXPERIENCE (CONTINUED)

TmaxOS & TmaxCloud

Gyeonggi, South Korea

Software & Platform Engineer

Jan. 2019 - Dec. 2020

- Built a Siri-like AI assistant on TmaxOS, covering model inference, data preprocessing, system integration, and asynchronous application flows.
- Designed and developed Kubernetes CI/CD and registry operators and deployed a private container image registry.

SELECTED ENGINEERING WORK

Internal AX Agent Platform

Built an AWS EKS platform for governed AX application delivery with runtime adaptation, automatic domains, authorized-only access, and integration with an existing enterprise LLM Gateway API.

Observability & Release Reliability

Deployed Grafana, Prometheus, and OpenTelemetry and connected Kubernetes releases across development, staging, and production with security and operational requirements.

Global B2C Generative AI Service

Supported UK rollout readiness, initial production operations, incident response, and early stabilization for a Microsoft-collaborative consumer AI service on Azure.

Enterprise AI Assistant | Public reference implementation

jinwonlee.github.io/projects/enterprise-ai-assistant/web/

Personal reference application combining SQLite FTS5 retrieval, role-based access filtering, grounded citations, audit events, FastAPI, and optional OpenAI Responses API integration.

EDUCATION & RESEARCH

UNIST | M.S. in Computer Science and Engineering

Mar. 2017 - Feb. 2019

Master's thesis: Convergence Aware CNN Training | Advisors: Jiwon Seo and Sam H. Noh

UNIST | B.S. in Computer Science and Engineering, Minor in Mathematics

Mar. 2012 - Feb. 2017

Undergraduate research: distributed systems for AI model training with TensorFlow

PUBLICATIONS & LANGUAGE

- USENIX NSDI 2019 poster: Alleviating the Network Bottleneck for CNN Distributed Training through Automatic Resource-Aware Layer Placement.
- arXiv 2019: Accelerated Training for CNN Distributed Deep Learning through Automatic Resource-Aware Layer Placement.
- ACM EuroSys 2018 poster: Improving Performance of Distributed TensorFlow using CNN Characteristics Exploiting Model Parallelism.
- OPIc IH (Intermediate High), Dec. 2025.