

Jinwon Lee

Software & AI Platform Engineer

Seoul, South Korea | +82 10-2422-1429 | dlwlsdnjsehs1993@gmail.com |
github.com/JINwonLEE | linkedin.com/in/jwl1993

"Build it first, then think it through."

TARGET FIT

- Application Software
- AI Platform
- End-to-End Delivery
- Reliability Engineering

SUMMARY

Software and AI platform engineer with about eight years of experience designing, developing, launching, and operating enterprise and user-facing applications. At Samsung Electronics, I build an internal AX Agent Platform and have delivered Kubernetes-based services, cloud automation, CI/CD, observability, and AI application workflows.

I translate application, security, and platform requirements into reliable software and carry delivery from architecture and implementation through rollout and production stabilization. I work across application, cloud, security, and operations concerns so teams can release services that users can depend on.

Application & Platform Software

End-to-End Product Delivery

Cloud-Native Engineering

Reliability & Performance

CORE SKILLS

Application Software	Enterprise and user-facing applications, Kubernetes operators, Python automation, asynchronous architecture, API integration
Product Delivery	Requirements translation, architecture, implementation, rollout readiness, release automation, production stabilization
Cloud Platform	Kubernetes, AWS EKS, Azure, OpenShift, Docker, Helm/Helmfile, ArgoCD, GitOps, Terraform, Ansible
Reliability	Authentication/authorization, security governance, domain provisioning, Grafana, Prometheus, OpenTelemetry, incident response

EXPERIENCE

Samsung Electronics (AX Development Group)

Seoul, South Korea

SRE / AI Platform Engineer

Jan. 2026 - Present

- Global B2C Generative AI Service (Jan. 2026 - Jun. 2026): Collaborated with Microsoft on UK rollout readiness, initial production operations, incident response, and early stabilization for an Azure-based consumer service.
- AX Agent Platform (Jul. 2026 - Present): Own end-to-end design and delivery of an internal platform on AWS EKS, adapting the agent runtime and deployment environment to enterprise platform standards.
- Translate deployment, authentication, authorization, security, and governance requirements into a self-service workflow with automatic domain provisioning and authorized-only access.
- Enabled deployed AX applications to consume the existing enterprise LLM Gateway API through platform-side client compatibility and integration.

Samsung Research

Seoul, South Korea

SRE / AI Platform Engineer

Jul. 2024 - Dec. 2025

- Led platform and application reliability for a company-wide AI chatbot service used by 30,000 employees on internal Kubernetes clusters.
- Owned production delivery across development, staging, and production environments, connecting application releases with security and operational requirements.
- Provisioned and operated development, staging, and production Kubernetes clusters using Helm and Helmfile.
- Built automated CI/CD pipelines for containerized applications with Docker, Kubernetes, GitHub Actions, and ArgoCD.
- Deployed centralized monitoring with Grafana, Prometheus, and OpenTelemetry, and built an automated LLM evaluation pipeline for Machine Translation tasks.

EXPERIENCE (CONTINUED)

● Samsung Electronics (Network Dept.)

Seoul, South Korea
Jan. 2021 - Jun. 2024

SRE & Software Engineer

- Designed and delivered containerized applications using Docker, Helm, Kubernetes Operator, and AWS services including EKS, ECS, Route 53, RDS, and IAM.
- Built and managed OpenShift, Kubernetes, and EKS environments for performance, reliability, high availability, and auto-scaling.
- Designed and developed an automated one-click Kubernetes cluster provisioning system for cloud developers.
- Developed an application that mirrors up to 130,000 packets per second in real time under constrained system resources.

● TmaxOS & TmaxCloud

Gyeonggi, South Korea
Jan. 2019 - Dec. 2020

Software & Platform Engineer

- Designed and developed CI/CD and registry operators in Kubernetes and deployed a private container image registry.
- Built a Siri-like AI assistant on TmaxOS with system-level integration, model inference, data preprocessing, and asynchronous application architecture.

EDUCATION

UNIST (Ulsan National Institute of Science and Technology)

Mar. 2017 - Feb. 2019

M.S. in Computer Science and Engineering

- Master Thesis: Convergence Aware CNN Training
- Advisors: Jiwon Seo (Hanyang University), Sam H. Noh (UNIST)

UNIST (Ulsan National Institute of Science and Technology)

Mar. 2012 - Feb. 2017

B.S. in Computer Science and Engineering; Minor in Mathematics

- Undergraduate research: distributed systems for AI model training with TensorFlow

CERTIFICATIONS & LANGUAGE

- OPIc IH (Intermediate High), Dec. 2025

PUBLICATIONS

- Alleviating the Network Bottleneck for CNN Distributed Training through Automatic Resource-Aware Layer Placement, USENIX NSDI 2019 poster.
- Accelerated Training for CNN Distributed Deep Learning through Automatic Resource-Aware Layer Placement, arXiv 2019.
- Improving Performance of Distributed TensorFlow using CNN Characteristics Exploiting Model Parallelism, ACM EuroSys 2018 poster.