

Jinwon Lee

AI Platform Software & SRE Engineer

Seoul, South Korea | +82 10-2422-1429 | dlwlsdnjsehs1993@gmail.com |
github.com/JINwonLEE | linkedin.com/in/jwl1993

"Build it first, then think it through."

TARGET FIT

- AI Platform Software
- MLOps / Model Serving
- Kubernetes Infrastructure
- System Software

SUMMARY

At Samsung Electronics, I design and build an internal AX Agent Platform while working on SRE and platform engineering for generative AI services. For about eight years, I have built and operated Kubernetes-based services, cloud infrastructure automation, CI/CD, observability, and AI evaluation/serving workflows.

My strength is taking AI applications beyond "working code" and turning them into systems that real users can rely on. I consider deployment, authentication, monitoring, incident response, and resource constraints together to build platforms where applications can run smoothly and reliably.

AI Platform Software

MLOps & Model Serving

System Software & Performance

Cloud/Kubernetes Infrastructure

CORE SKILLS

AI/MLOps	LLM evaluation pipeline, MLflow metric tracking, AI serving operations, Document LLM platform, Langfuse observability
Platform	Kubernetes, AWS EKS, Azure, OpenShift, Docker, Helm/Helmfile, ArgoCD, GitOps, Terraform, Ansible
System	Linux runtime behavior, Python automation, high-throughput packet mirroring, distributed systems, resource-aware optimization
Reliability	Authentication/authorization, security governance, domain provisioning, Grafana, Prometheus, OpenTelemetry, incident response

EXPERIENCE

Samsung Electronics (AX Development Group)

Seoul, South Korea

SRE / AI Platform Engineer

Jan. 2026 - Present

- Global B2C Generative AI Service (Jan. 2026 - Jun. 2026): Supported the UK expansion, initial production operations, and incident response for a Microsoft Azure-based service in collaboration with Microsoft.
- AX Agent Platform (Jul. 2026 - Present): Designed and built an internal platform on AWS EKS, adapting the agent runtime and deployment environment to internal platform standards.
- Owned authentication, authorization, security, and internal governance while building self-service deployment with automatic domain provisioning and authorized-only access.
- Enabled deployed AX applications to consume the existing enterprise LLM Gateway API through platform-side client compatibility and integration, without owning the Gateway design.

Samsung Research

Seoul, South Korea

SRE / AI Platform Engineer

Jul. 2024 - Dec. 2025

- Operated a company-wide AI chatbot service on internal Kubernetes clusters, owning platform and application reliability for production use.
- Led infrastructure, platform, and application operations for a company-wide AI chatbot platform serving 30,000 employees on Kubernetes.
- Provisioned and operated development, staging, and production Kubernetes clusters using Helm and Helmfile.
- Built automated CI/CD pipelines for containerized applications with Docker, Kubernetes, GitHub Actions, and ArgoCD.
- Deployed centralized monitoring with Grafana, Prometheus, and OpenTelemetry, and built an automated LLM evaluation pipeline for Machine Translation tasks.

EXPERIENCE (CONTINUED)

Samsung Electronics (Network Dept.)

Seoul, South Korea
Jan. 2021 - Jun. 2024

SRE & Software Engineer

- Deployed containerized applications using Docker, Helm, Kubernetes Operator, and AWS services including EKS, ECS, Route 53, RDS, and IAM.
- Built and managed OpenShift, Kubernetes, and EKS environments with focus on performance, reliability, high availability, and auto-scaling.
- Developed an automated one-click Kubernetes cluster provisioning system for cloud developers.
- Developed an application that mirrors up to 130,000 packets per second in real time under constrained system resources.

TmaxOS & TmaxCloud

Gyeonggi, South Korea
Jan. 2019 - Dec. 2020

Software & Platform Engineer

- Developed CI/CD and registry operators in Kubernetes and deployed a private container image registry.
- Built a Siri-like AI assistant on TmaxOS with system-level integration, AI model inference, preprocessing, and asynchronous architecture.

EDUCATION

UNIST (Ulsan National Institute of Science and Technology)

Mar. 2017 - Feb. 2019

M.S. in Computer Science and Engineering

- Master Thesis: Convergence Aware CNN Training
- Advisors: Jiwon Seo (Hanyang University), Sam H. Noh (UNIST)

UNIST (Ulsan National Institute of Science and Technology)

Mar. 2012 - Feb. 2017

B.S. in Computer Science and Engineering; Minor in Mathematics

- Undergraduate research: distributed systems for AI model training with TensorFlow

CERTIFICATIONS & LANGUAGE

- OPIc IH (Intermediate High), Dec. 2025

PUBLICATIONS

- Alleviating the Network Bottleneck for CNN Distributed Training through Automatic Resource-Aware Layer Placement, USENIX NSDI 2019 poster.
- Accelerated Training for CNN Distributed Deep Learning through Automatic Resource-Aware Layer Placement, arXiv 2019.
- Improving Performance of Distributed TensorFlow using CNN Characteristics Exploiting Model Parallelism, ACM EuroSys 2018 poster.